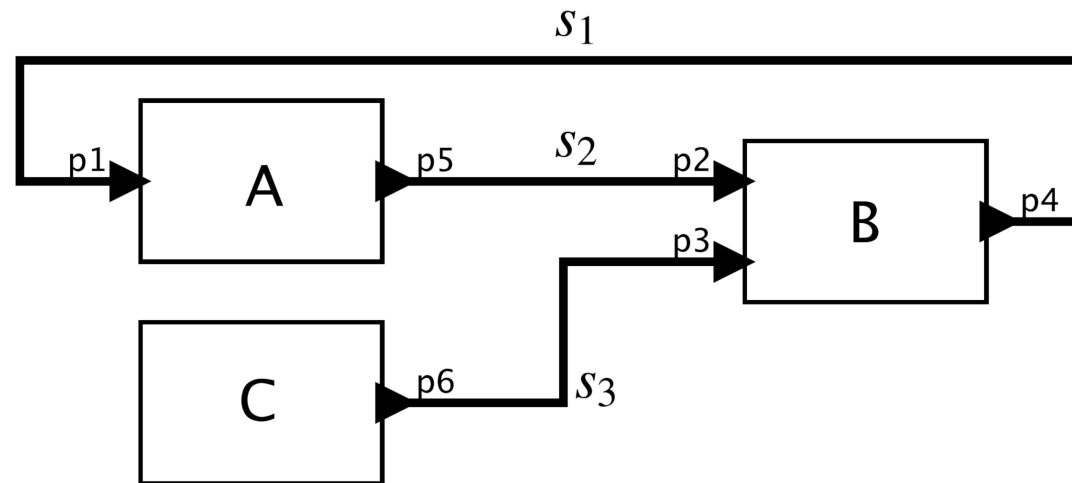


Kahn Process Networks

- A set of computational components, called *actors*.
- Each representing a sequential procedure (process unit).
 - receiving messages from other actors
 - performing local operations
 - sending messages to other actors
- Messages are communicated asynchronously with unbounded buffers.
- A procedure can always send a message. It does not need to wait for the recipient to be ready to receive.
- Messages are delivered reliably and in order.
- When a procedure attempts to receive a message, that attempt blocks the procedure until a message is available.

Example 1



Example of a process network

- Kahn process networks
 - least fixed point semantics, deterministic
 - unbounded buffer
 - grid databases: one request/one response
- No unbounded buffer

Desired Properties of concurrent, parallel, distributed, event-driven systems

- asynchronism
- determinism
- dead lock free
- scalability

MapReduce

- Model for processing large data sets
- Contains Map and Reduce functions
- Runs on a large cluster of machines
- A lot of MapReduce programs are executed on Google's cluster everyday

Motivation of MapReduce

- Input data is large
 - the whole web, billions of pages
- Lots of machines
 - how to use them efficiently

Programming model

- Input and output
 - each a set of key/value pairs
- Programmer specifies two functions
 - $\text{map}(\text{in-key}, \text{in-value}) \longrightarrow \text{list}(\text{out-key}, \text{intermediate-value})$
 - * processing input key/value pair
 - * process set of intermediate pairs
 - $\text{reduce}(\text{out-key}, \text{list}(\text{intermediate-value})) \longrightarrow \text{list}(\text{out-value})$
 - * Combines all intermediate values for a particular key
 - * Produces a set of merged output values (usually just one)

Example Word counter

Input: a document with 3 pages

- Page 1: the weather is good
- Page 2: today is good
- Page 3: good weather is good

Map Output

- Worker 1 (one processor)
 - (the 1), (weather 1), (is 1), (good 1)
- Worker 2
 - (today 1), (is 1), (good 1)
- Worker 3
 - (good 1), (weather 1), (is 1), (good 1)

Reduce Input

- Worker 1
 - (the 1),
- Worker 2
 - (is 1), (is 1), (is 1)
- Worker 3
 - (weather 1), (weather 1)
- Worker 4
 - (today 1)
- Worker 5
 - (good 1), (good 1), (good 1), (good 1)

Reduce Output

- Worker 1
 - (the 1),
- Worker 2
 - (is 3)
- Worker 3
 - (weather 2)
- Worker 4
 - (today 1)
- Worker 5
 - (good 4),

Map Input

- Page 1
the weather is good
- Page 2
today is good
- Page 3
good weather is good

Map Output

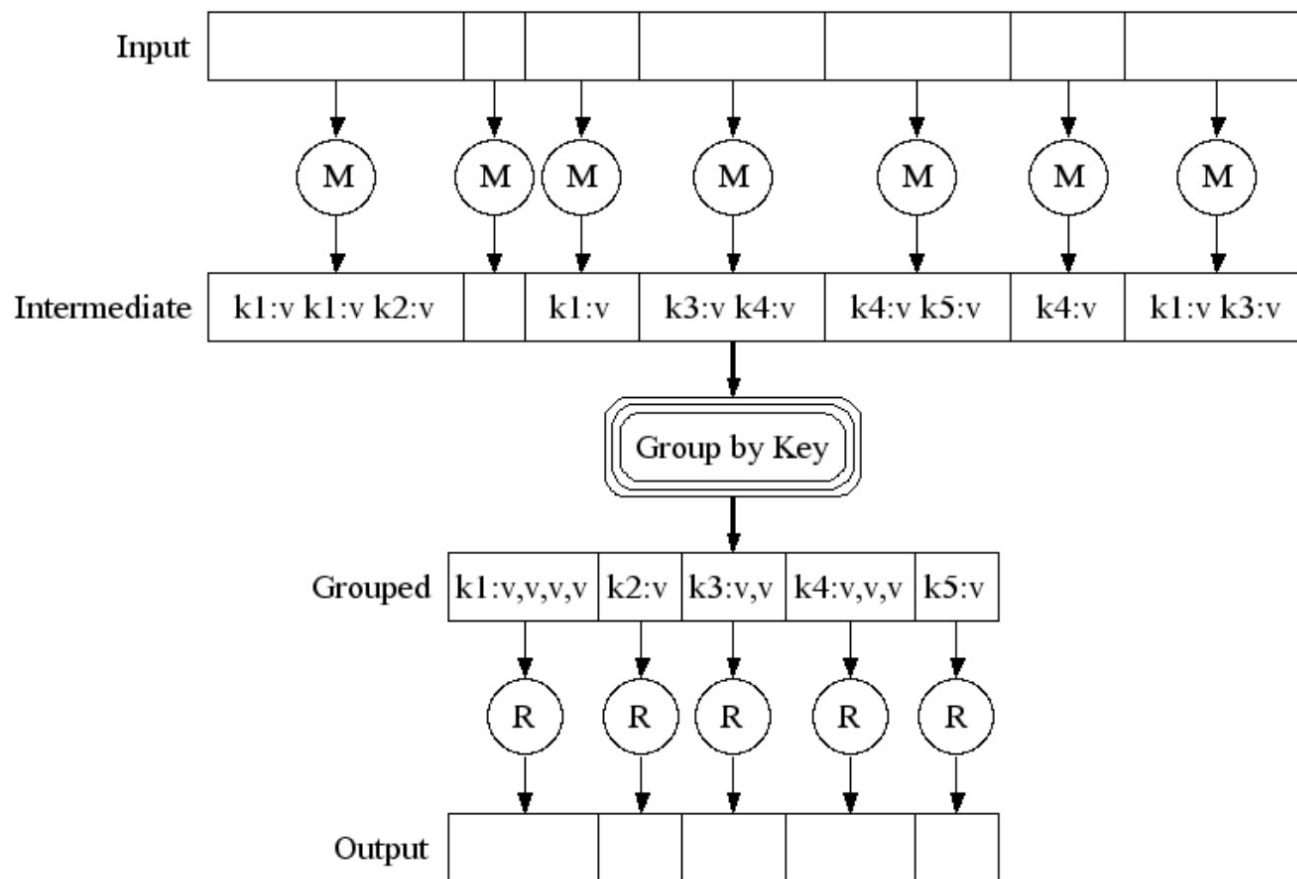
- Worker 1 (one processor)
 - (the 1), (weather 1), (is 1), (good 1)
- Worker 2
 - (today 1), (is 1), (good 1)
- Worker 3
 - (good 1), (weather 1), (is 1), (good 1)

Reduce Input

- Worker 1
 - (the 1),
- Worker 2
 - (is 1), (is 1), (is 1)
- Worker 3
 - (weather 1), (weather 1)
- Worker 4
 - (today 1)
- Worker 5
 - (good 1), (good 1), (good 1), (good 1)

Reduce Output

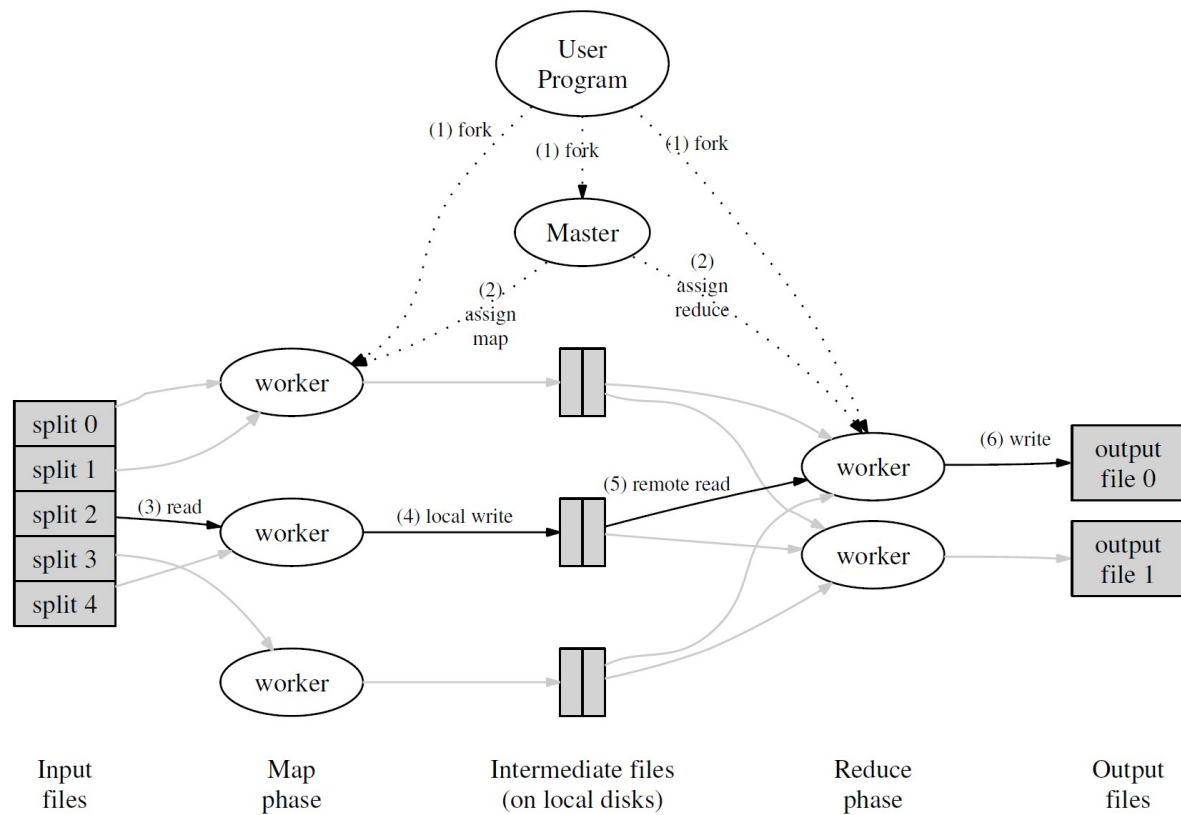
- Worker 1
 - (the 1),
- Worker 2
 - (is 3)
- Worker 3
 - (weather 2)
- Worker 4
 - (today 1)
- Worker 5
 - (good 4),




Other Examples

- Inverted Index
 - Input: list of $\langle \text{docId}, \text{document} \rangle$
 - Output: list of $\langle \text{Word}, \text{list}(\text{docId}, \text{frequency}) \rangle$
 - Map: parse each document and outputs a sequence of $\langle \text{word}, \text{docId} \rangle$
 - Reduce: accept all pairs for a given word, sorts docId according its frequency, and generate the output
- Count of URL access frequency
 - Map function
process logs of web page requests and outputs $\langle \text{URL}, 1 \rangle$
 - Reduce function
add all values for the same URL and outputs a $\langle \text{URL}, \text{total count} \rangle$

Execution Overview




MapReduce vs Kahn Process Network



Implementation of Kahn process networks

- KPN scheduling



- Message transport
- Deadlock detection and resolution

Experimental

Word count Count the number of occurrences of each word in a given text document and outputs them sorted in the order of decreasing frequency.

Experimental

- File size: 10, 50, and 100 (MBs)
- Machine: a single server with 8 CPUs

Three solutions

- Phoenix


A MapReduce implementation designed specifically for multi-core machines.

- KPN-MR

A Kahn process network configured for MapReduce using two instantiations connected in series

$$S \longrightarrow MR_1 \longrightarrow MR_2 \longrightarrow M$$

- KPN-FLEX



Results

Doc Size (MB)	Phoenix (s)	KPN-MR (s)	KPN-FLEX (s)
10	0.30	0.32	0.11
50	0.98	1.86	0.37
100	2.04	4.23	0.74

Experimental for K-means

Points/Group	Phoenix (s)	KPN-MR (s)	KPN-FLEX (s)
100K/100	2.39	1.83	1.51
200K/50	3.92	3.20	2.24
100K/100	5.43	5.19	4/26